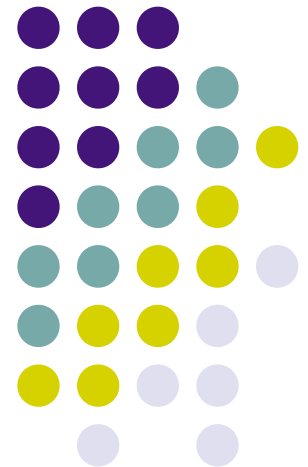


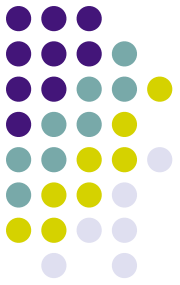
# Chord

---

A scalable peer-to-peer  
look-up protocol for  
internet applications



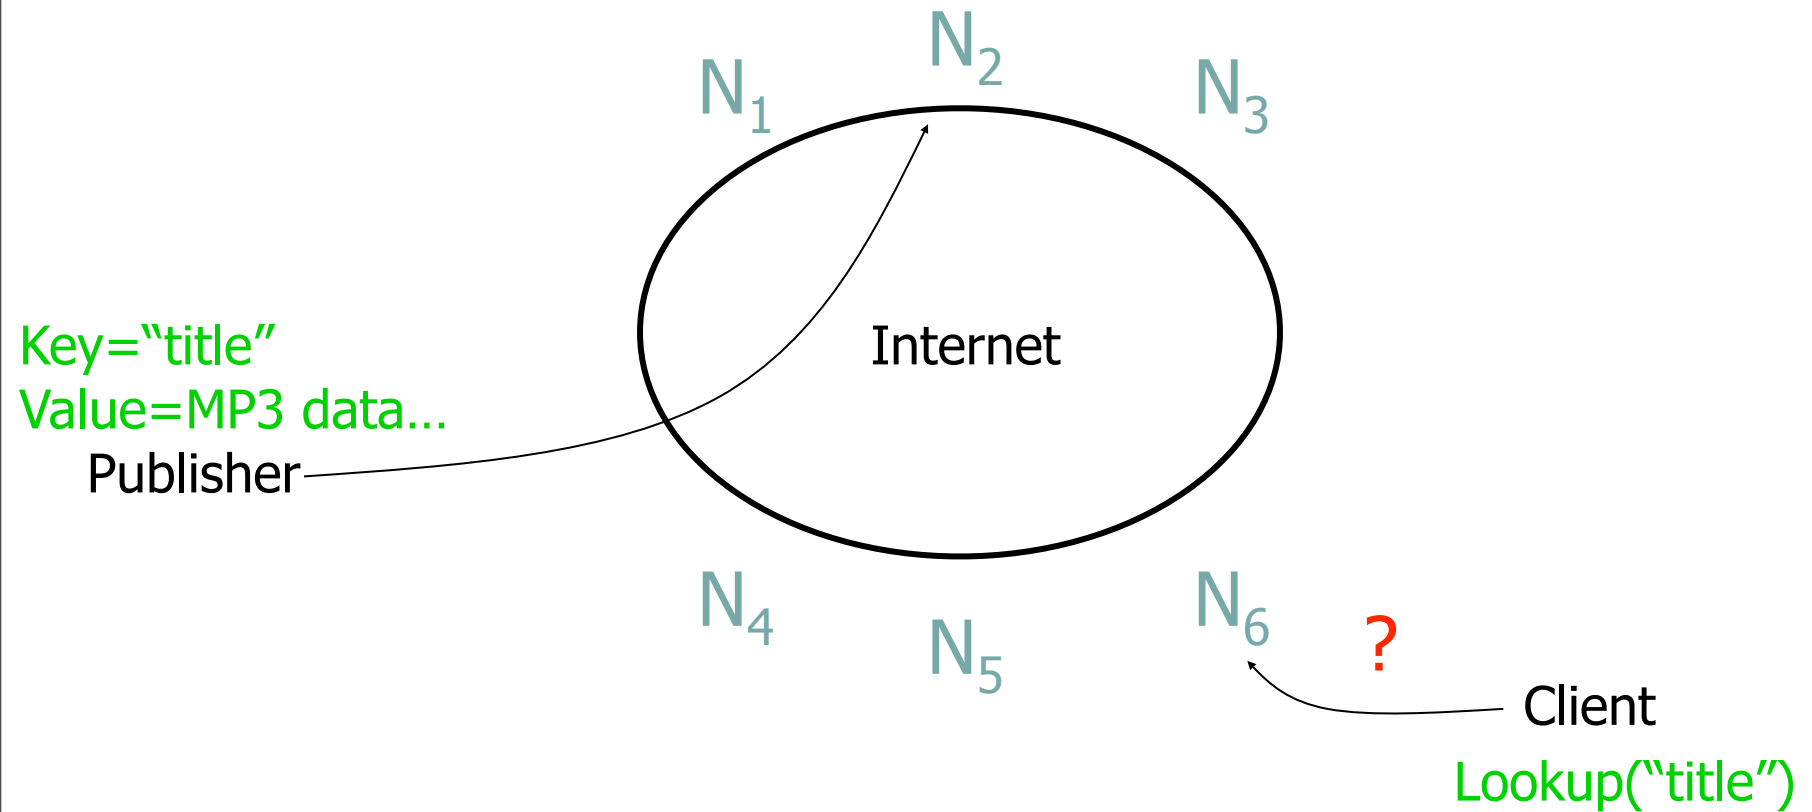
by Ion Stoica, Robert Morris, David Karger,  
M. Frans Kaashoek, Hari Balakrishnan

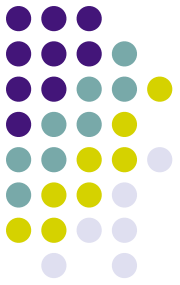


# Overview


- Introduction
- The Chord Algorithm
  - Construction of the Chord ring
  - Localization of nodes
  - Node joins and stabilization
  - Failure of nodes
- Applications
- Summary
- Questions

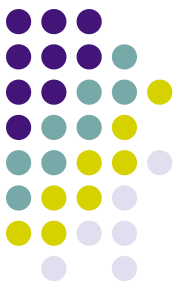
# The lookup problem





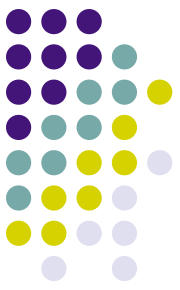
# Routed queries (Freenet, Chord, etc.)





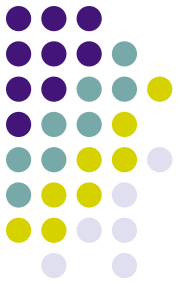
# What is Chord?

- Problem addressed: efficient node localization
- Distributed lookup protocol
- Simplicity, provable performance, proven correctness
- Support of just one operation: given a key, Chord maps the key onto a node



# Chord software

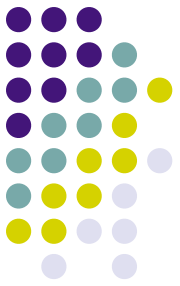
- 3000 lines of C++ code
- Library to be linked with the application
- provides a `lookup(key)` – function: yields the IP address of the node responsible for the key
- Notifies the node of changes in the set of keys the node is responsible for



# Overview

- Introduction
- The Chord Algorithm
  - Construction of the Chord ring
  - Localization of nodes
  - Node joins and Stabilization
  - Failure/Departure of nodes
- Applications
- Summary
- Questions

# The Chord algorithm – Construction of the Chord ring



- use Consistent Hash Function assigns each node and each key an m-bit identifier using SHA 1 (Secure Hash Standard).

m = any number big enough to make collisions improbable

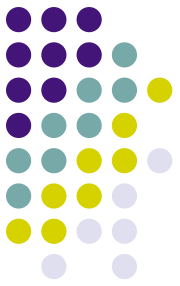
Key identifier = SHA-1(key)

Node identifier = SHA-1(IP address)


- Both are uniformly distributed
- Both exist in the same ID space



# The Chord algorithm – Construction of the Chord ring




- identifiers are arranged on a identifier circle modulo  $2^m \Rightarrow$  **Chord ring**



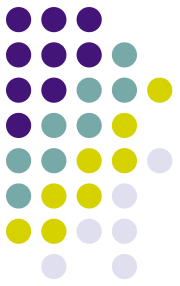
# The Chord algorithm – Construction of the Chord ring



- a key  $k$  is assigned to the node whose identifier is equal to or greater than the key's identifier
- this node is called  $\text{successor}(k)$  and is the first node clockwise from  $k$ .




# The Chord algorithm – Simple node localization

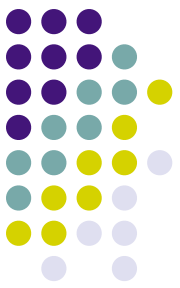


```
// ask node n to find the successor of id
n.find_successor(id)
  if (id ∈ (n; successor])
    return successor;
  else
    // forward the query around the
    // circle
    return successor.find_successor(id);
```

**=> Number of messages linear in  
the number of nodes !**

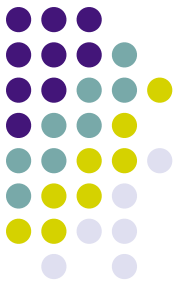


# The Chord algorithm – Scalable node localization



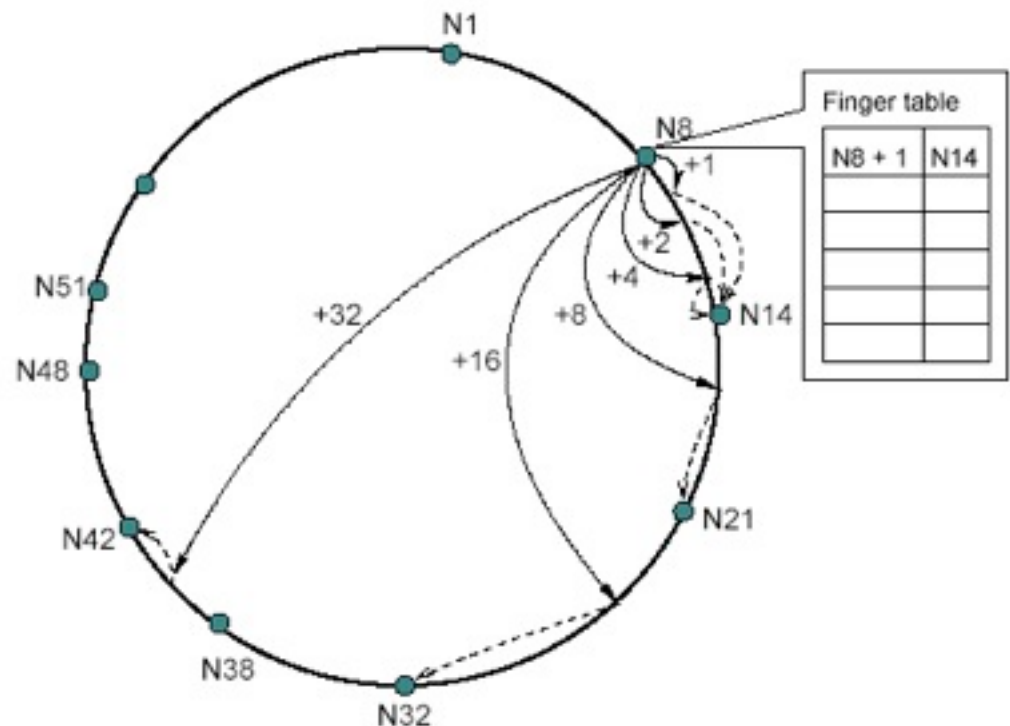
- Additional routing information to accelerate lookups
- Each node  $n$  contains a routing table with up to  $m$  entries ( $m$ : number of bits of the identifiers)  $\Rightarrow$  finger table
- $i^{\text{th}}$  entry in the table at node  $n$  contains the first node  $s$  that succeeds  $n$  by at least  $2^{i-1}$
- $s = \text{successor}(n + 2^{i-1})$
- $s$  is called the  $i^{\text{th}}$  finger of node  $n$

# The Chord algorithm – Scalable node localization

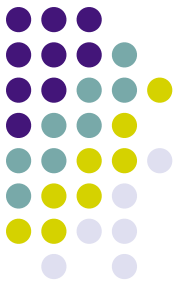


**Finger table:**

$finger[i] =$   
 $successor(n + 2^{i-1})$

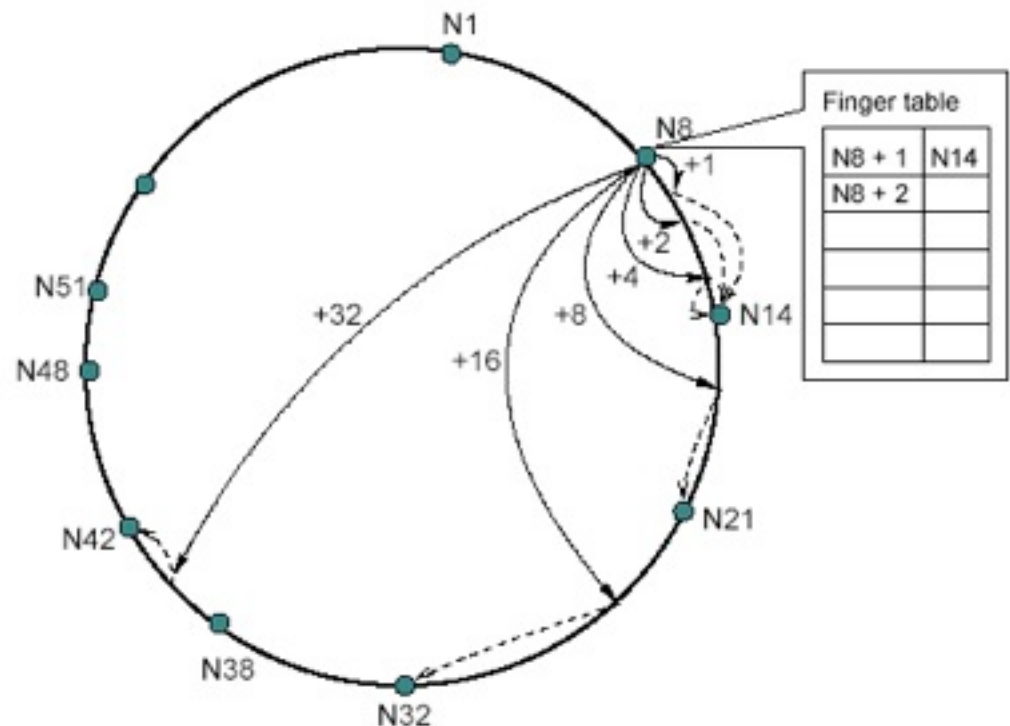


# The Chord algorithm – Scalable node localization

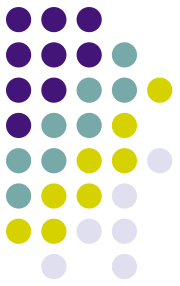


**Finger table:**

$finger[i] =$   
 $successor(n + 2^{i-1})$

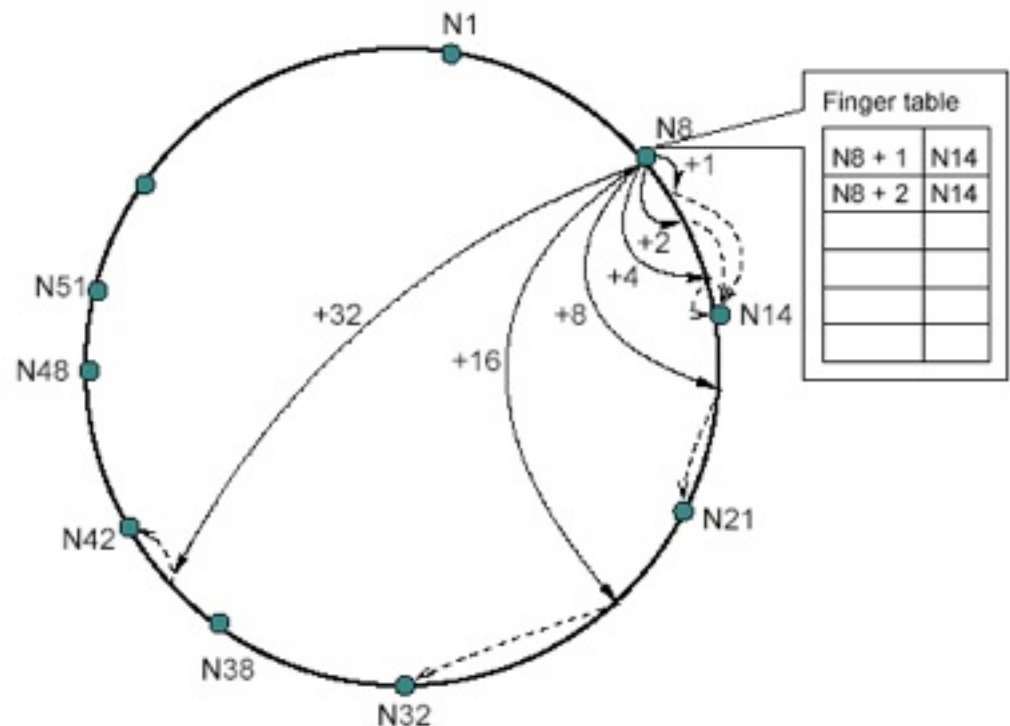


# The Chord algorithm – Scalable node localization

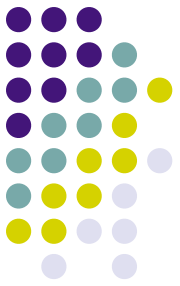


**Finger table:**

$finger[i] =$   
 $successor(n + 2^{i-1})$

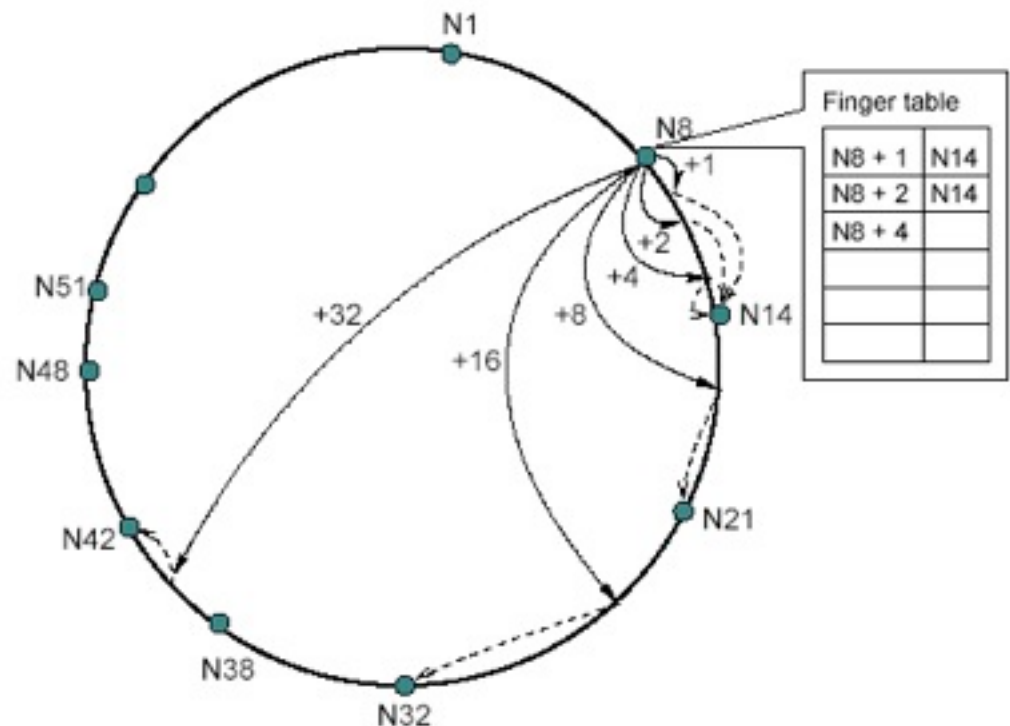


# The Chord algorithm – Scalable node localization



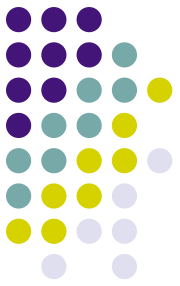
**Finger table:**

$finger[i] =$   
 $successor(n + 2^{i-1})$



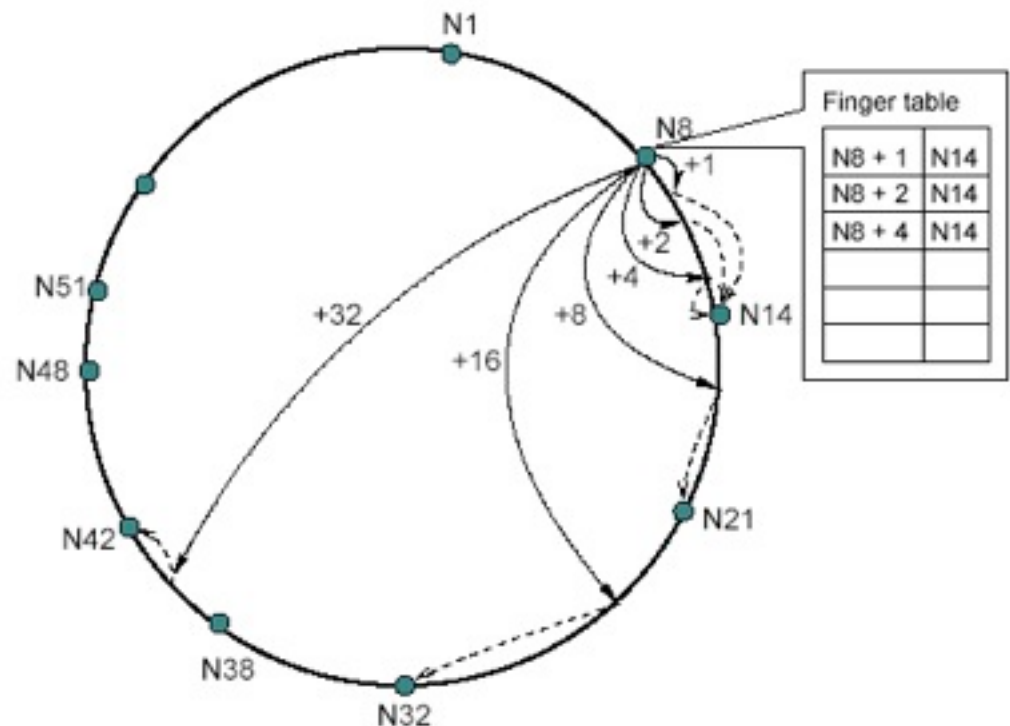


# The Chord algorithm – Scalable node localization



**Finger table:**

$finger[i] =$   
 $successor(n + 2^{i-1})$

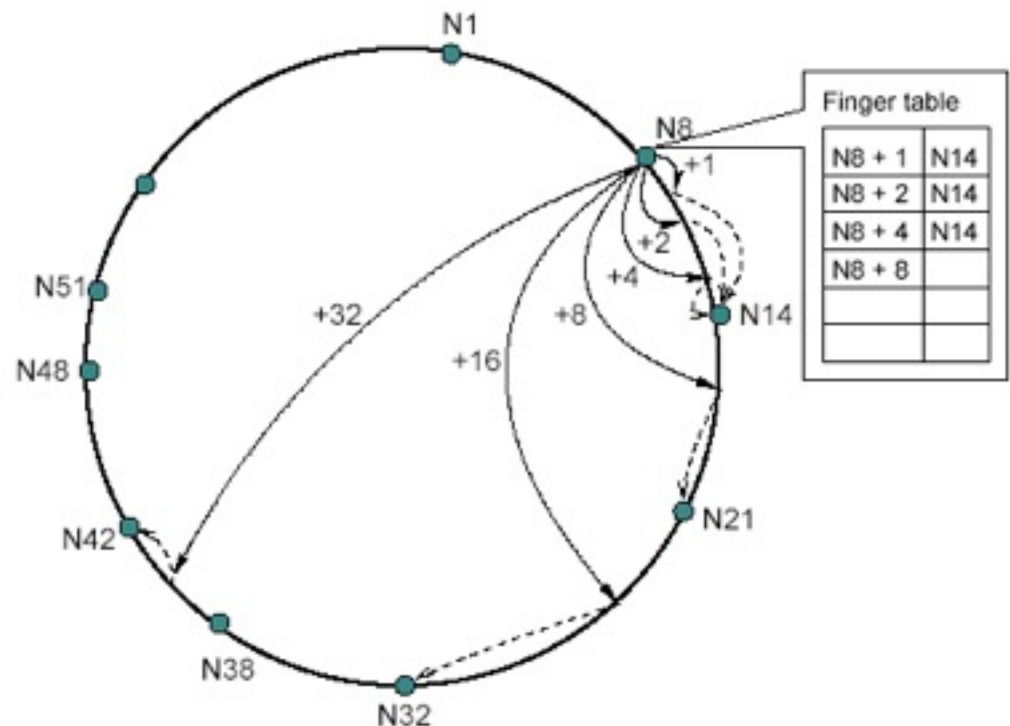


# The Chord algorithm – Scalable node localization

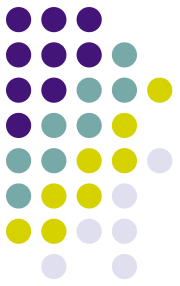


**Finger table:**

$finger[i] =$   
 $successor(n + 2^{i-1})$

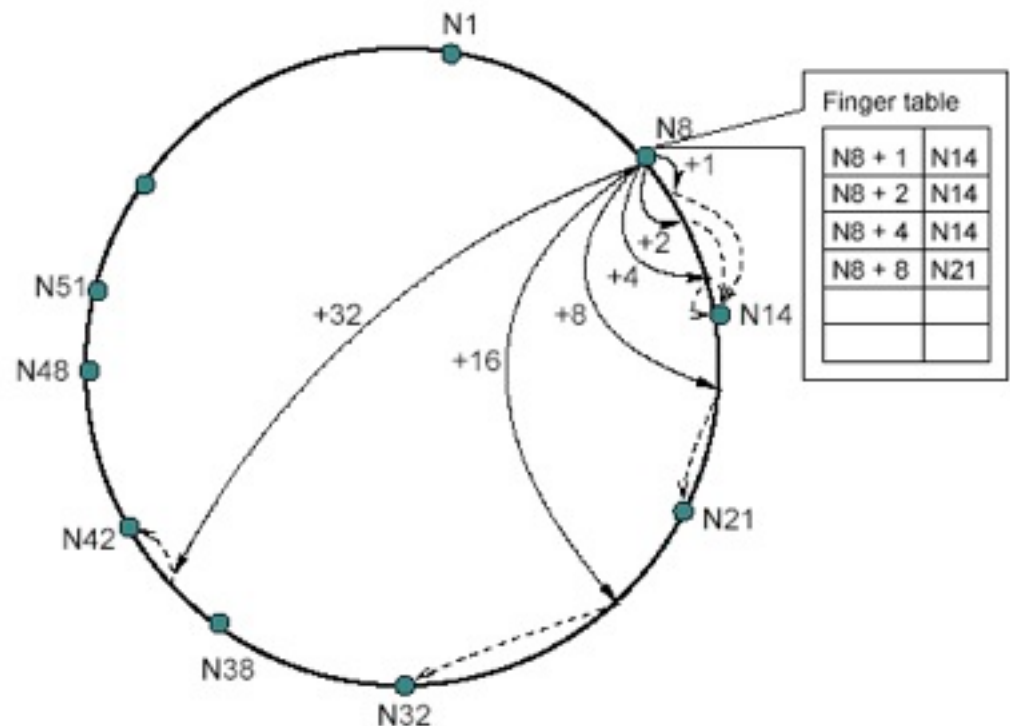


# The Chord algorithm – Scalable node localization

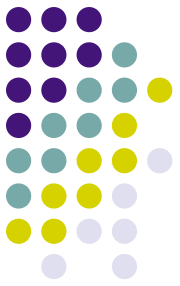


**Finger table:**

$finger[i] =$   
 $successor(n + 2^{i-1})$

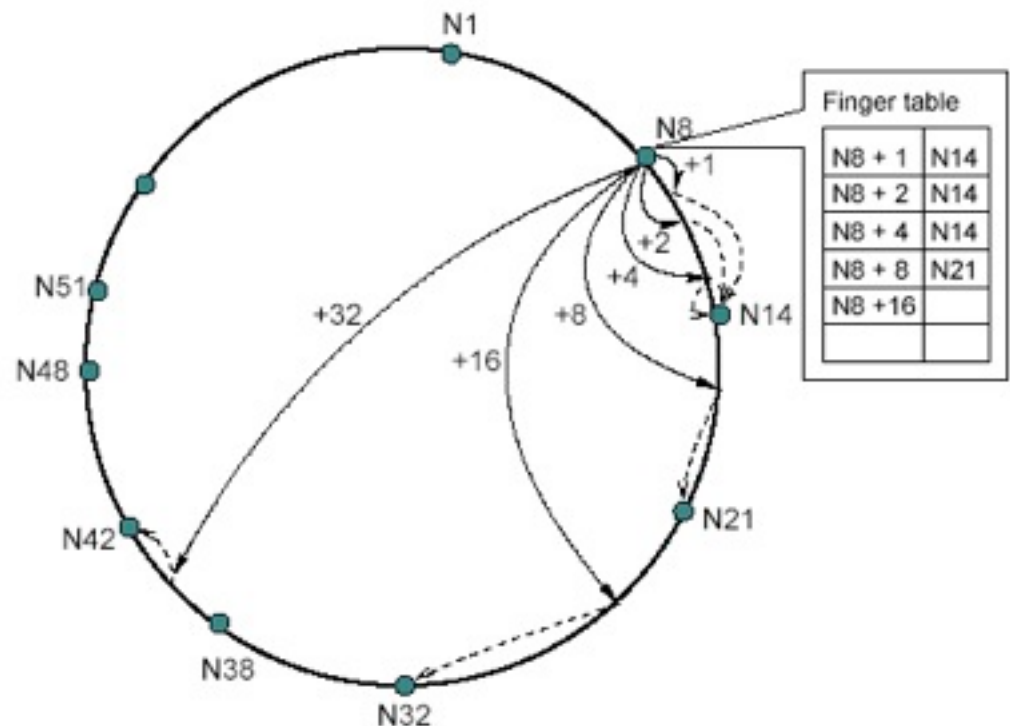


# The Chord algorithm – Scalable node localization

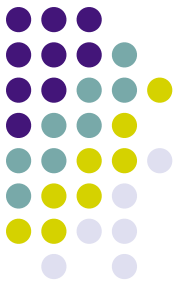


**Finger table:**

$finger[i] =$   
 $successor(n + 2^{i-1})$

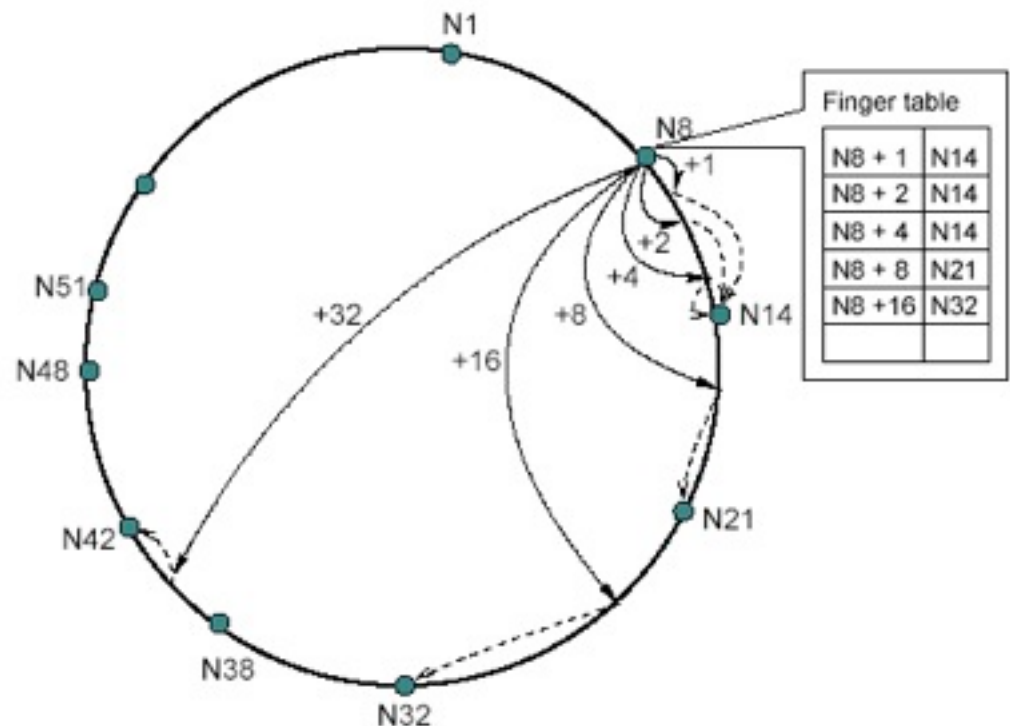


# The Chord algorithm – Scalable node localization

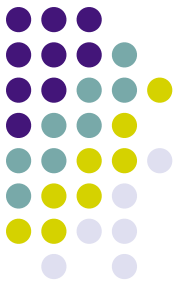


**Finger table:**

$finger[i] =$   
 $successor(n + 2^{i-1})$

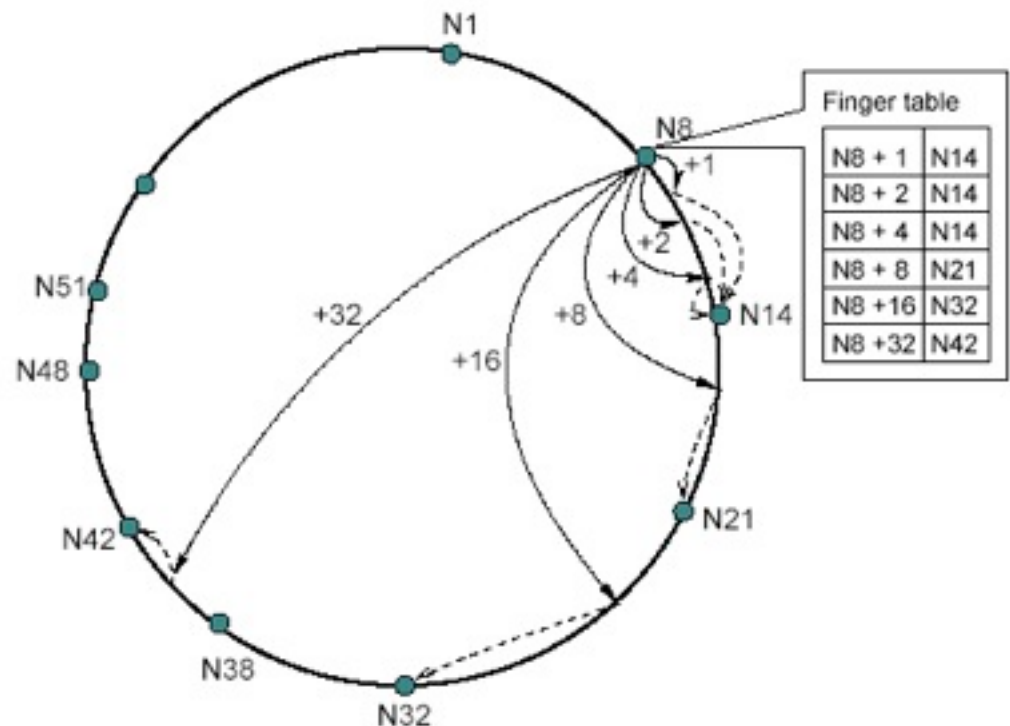


# The Chord algorithm – Scalable node localization

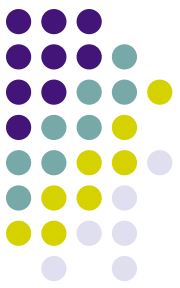


**Finger table:**

$finger[i] =$   
 $successor(n + 2^{i-1})$




# The Chord algorithm – Scalable node localization



Important characteristics of this scheme:


- Each node stores information about only a small number of nodes ( $m$ )
- Each nodes knows more about nodes closely following it than about nodes further away
- A finger table generally does not contain enough information to directly determine the successor of an arbitrary key  $k$

# The Chord algorithm – Scalable node localization



- Search in finger table for the nodes which most immediately precedes id
- Invoke `find_successor` from that node

it didn't forward the search to its successor

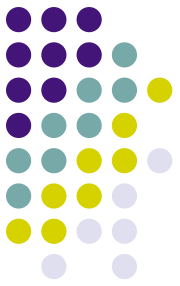


Finger table	
N8 + 1	N14
N8 + 2	N14
N8 + 4	N14
N8 + 8	N21
N8 + 16	N32
N8 + 32	N42

**=> Number of messages  $O(\log N)$ !**

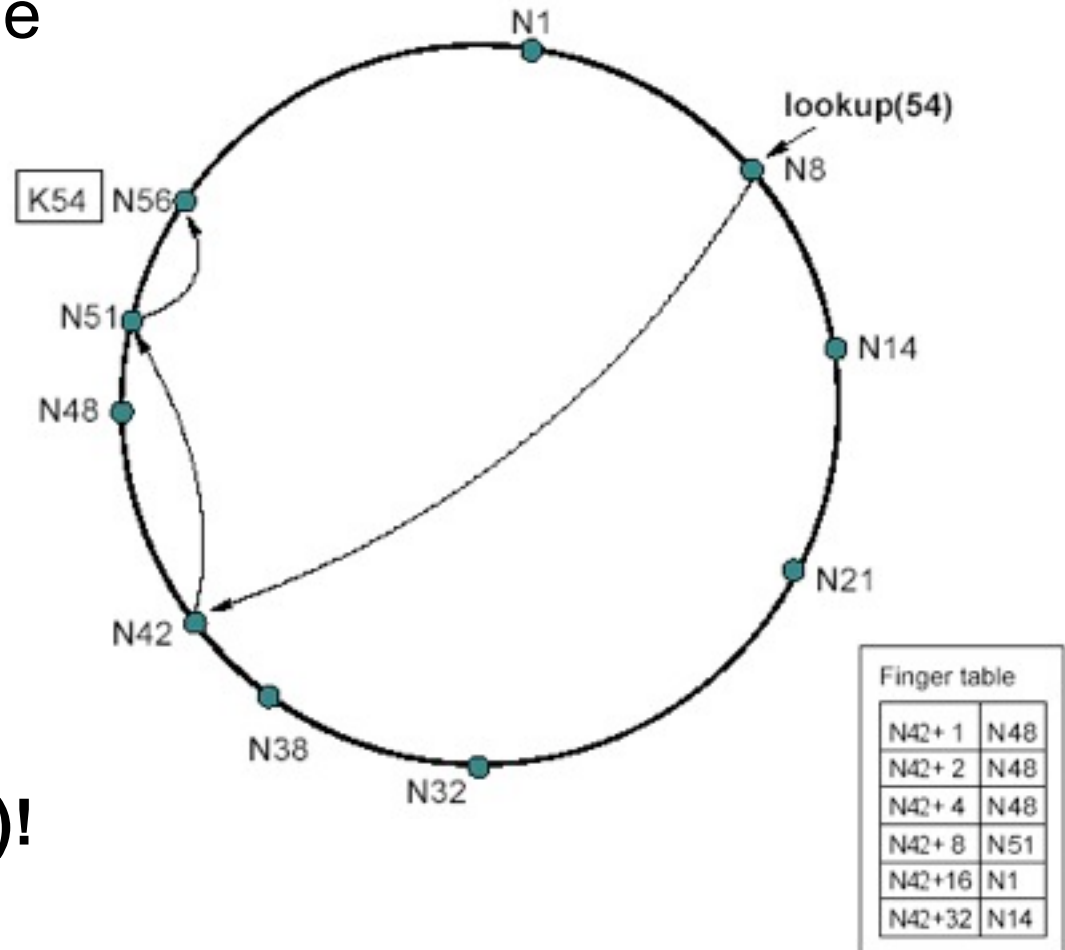


# The Chord algorithm – Scalable node localization

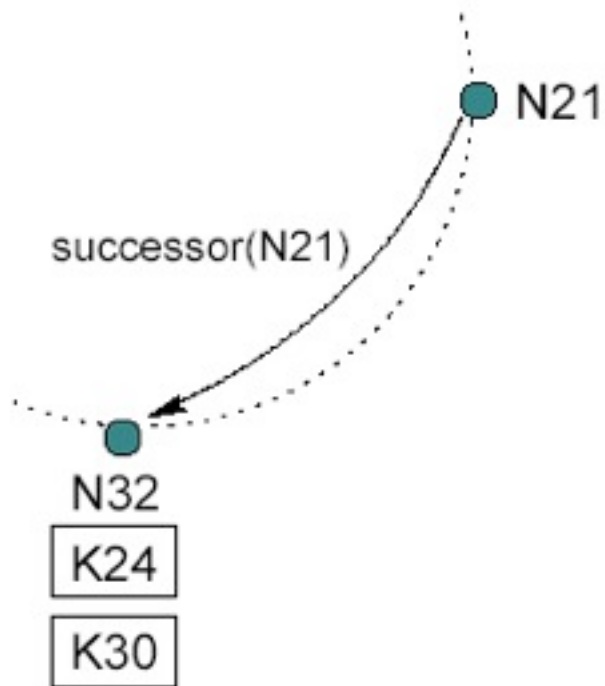


- Search in finger table for the nodes which most immediately precedes id
- Invoke `find_successor` from that node

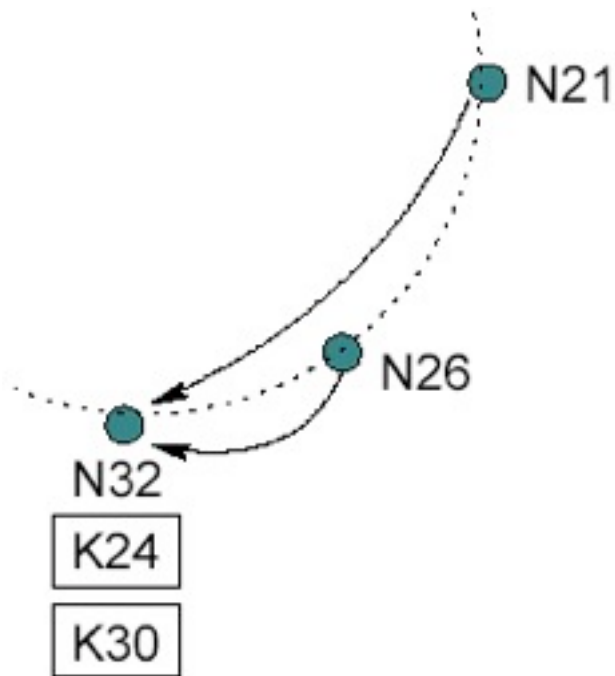
**=> Number of messages  $O(\log N)$ !**



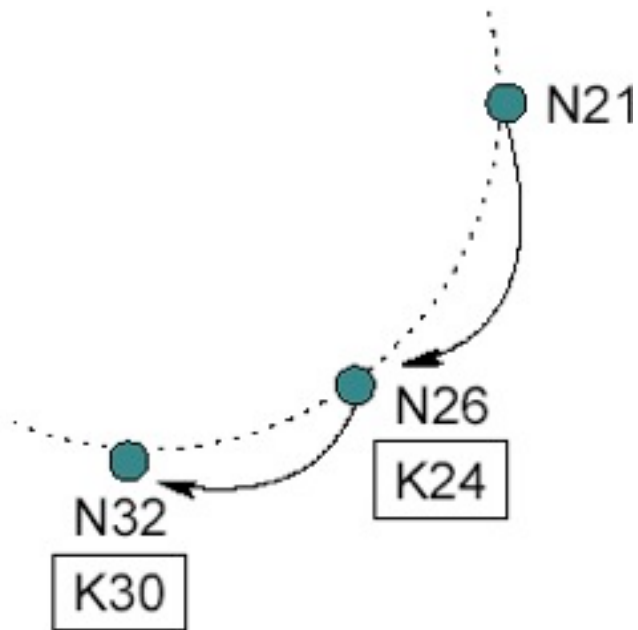
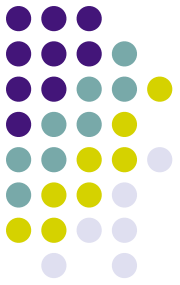
# The Chord algorithm – Node joins and stabilization



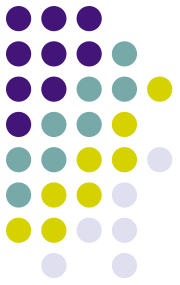
# The Chord algorithm – Node joins and stabilization



# The Chord algorithm – Node joins and stabilization



# The Chord algorithm – Node joins and stabilization



- To ensure correct lookups, all successor pointers must be up to date
- => stabilization protocol running periodically in the background
- Updates finger tables and successor pointers

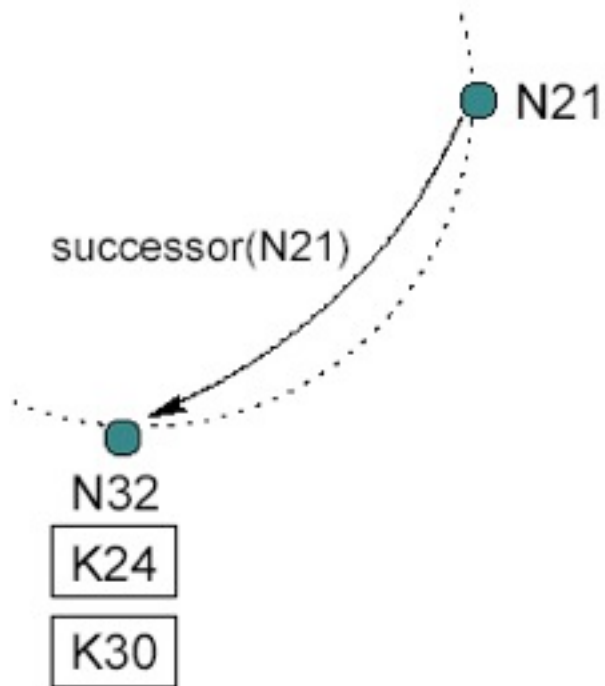
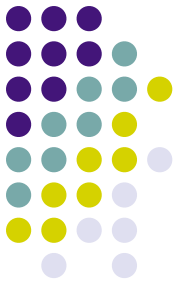
# The Chord algorithm – Node joins and stabilization



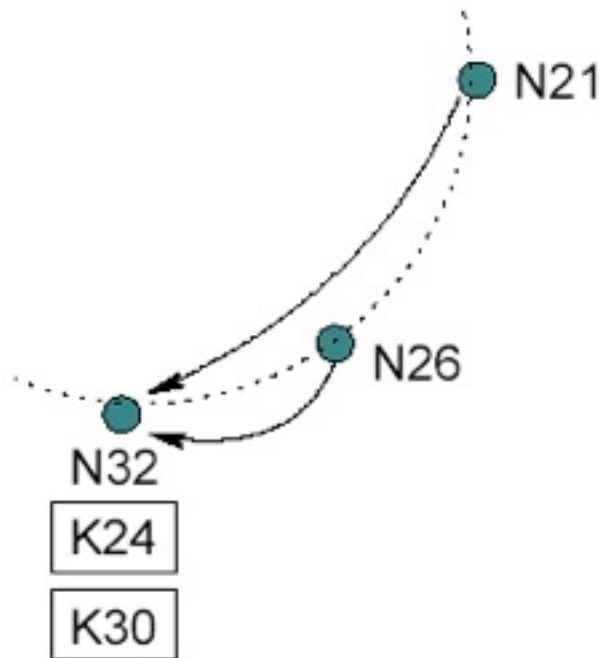
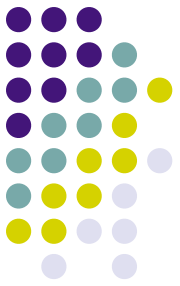
Stabilization protocol:

- **Stabilize():**  $n$  asks its successor for its predecessor  $p$  and decides whether  $p$  should be  $n$ 's successor instead (this is the case if  $p$  recently joined the system).
- **Notify():** notifies  $n$ 's successor of its existence, so it can change its predecessor to  $n$
- **Fix\_fingers():** updates finger tables

# The Chord algorithm – Node joins and stabilization




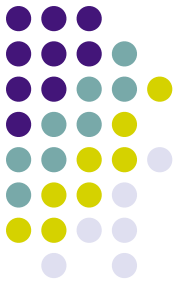
# The Chord algorithm – Node joins and stabilization



- N26 joins the system
- N26 acquires N32 as its successor
- N26 notifies N32
- N32 acquires N26 as its predecessor




# The Chord algorithm – Node joins and stabilization



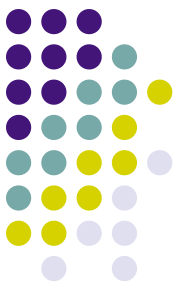
- N26 copies keys
- N21 runs stabilize() and asks its successor N32 for its predecessor which is N26.

# The Chord algorithm – Node joins and stabilization




- N21 acquires N26 as its successor
- N21 notifies N26 of its existence
- N26 acquires N21 as predecessor

# The Chord algorithm – Impact of node joins on lookups

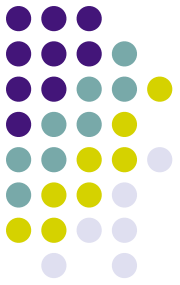


- All finger table entries are correct  $\Rightarrow O(\log N)$  lookups
- Successor pointers correct, but fingers inaccurate  $\Rightarrow$  correct but slower lookups

Finger table	
$N8 + 1$	N14
$N8 + 2$	N14
$N8 + 4$	N14
$N8 + 8$	N21
$N8 + 16$	N32
$N8 + 32$	N42

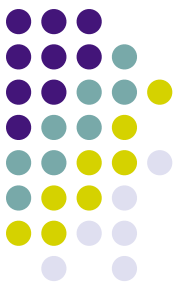


# The Chord algorithm – Impact of node joins on lookups



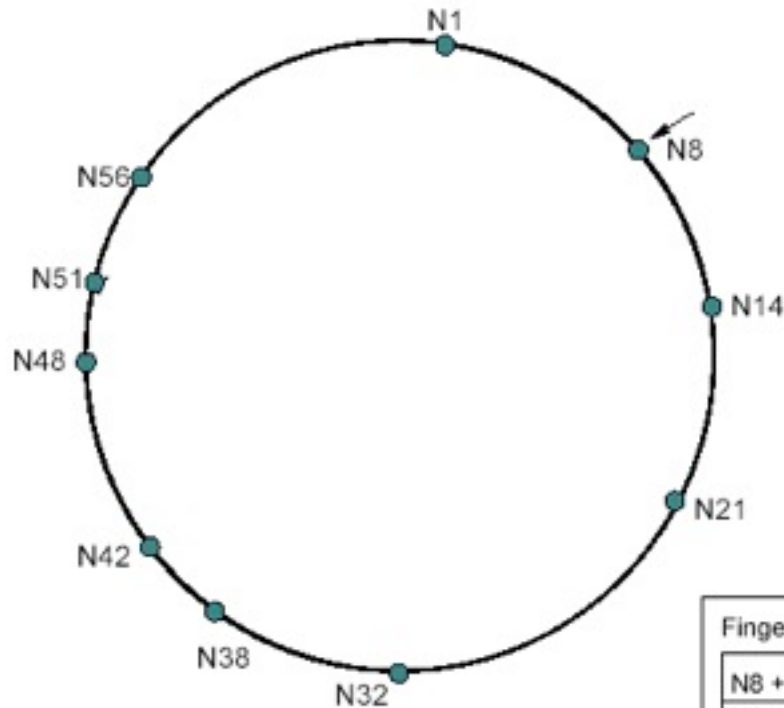
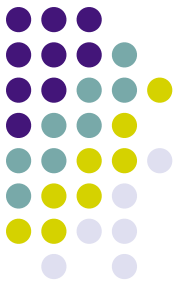
- Incorrect successor pointers => lookup might fail, retry after a pause
- But still correctness!

# The Chord algorithm – Impact of node joins on lookups



- Stabilization completed => no influence on performance
- Only for the negligible case that a large number of nodes joins between the target's predecessor and the target, the lookup is slightly slower
- No influence on performance as long as fingers are adjusted faster than the network doubles in size

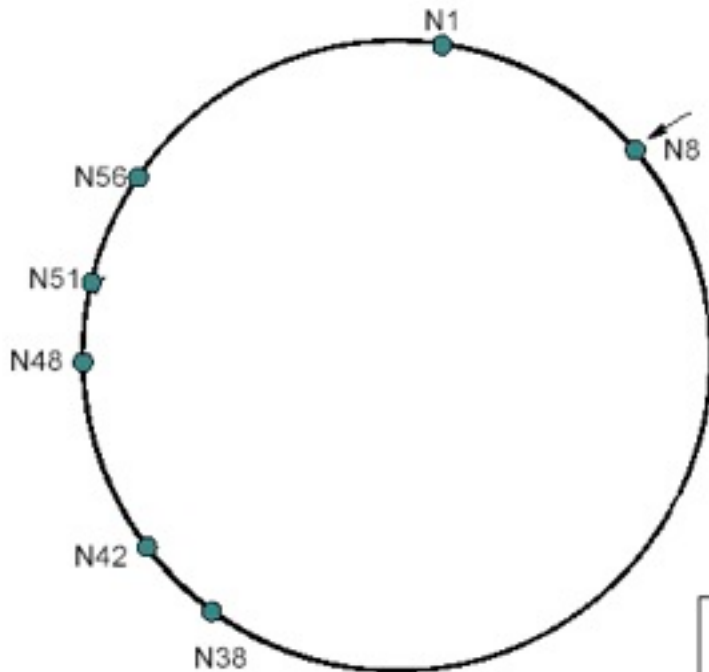
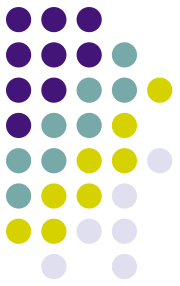
# The Chord algorithm – Failure of nodes



Finger table	
$N8 + 1$	N14
$N8 + 2$	N14
$N8 + 4$	N14
$N8 + 8$	N21
$N8 + 16$	N32
$N8 + 32$	N42

- Correctness relies on correct successor pointers
- What happens, if N14, N21, N32 fail simultaneously?
- How can N8 acquire N38 as successor?

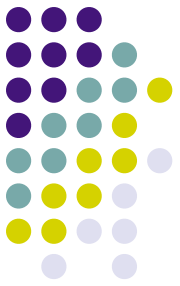
# The Chord algorithm – Failure of nodes



Finger table	
N8 + 1	N14
N8 + 2	N14
N8 + 4	N14
N8 + 8	N21
N8 + 16	N32
N8 + 32	N42

- Correctness relies on correct successor pointers
- What happens, if N14, N21, N32 fail simultaneously?
- How can N8 acquire N38 as successor?

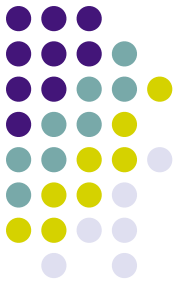
# The Chord algorithm – Failure of nodes



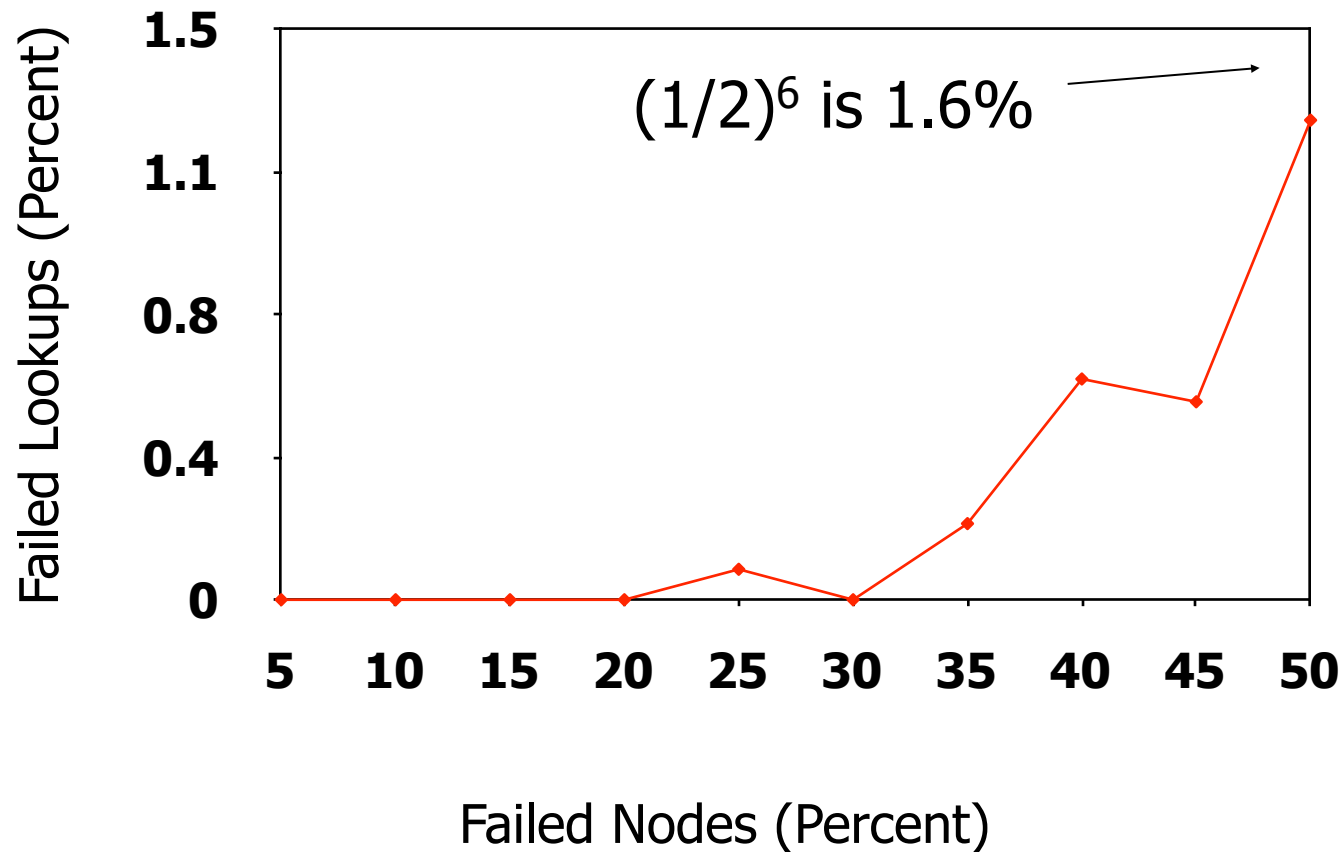
- Each node maintains a successor list of size  $r$
- If the network is initially stable, and every node fails with probability  $\frac{1}{2}$ , `find_successor` still finds the closest living successor to the query key and the expected time to execute `find_successor` is  $O(\log N)$
- Proofs are in the paper

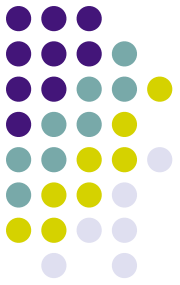


# The Chord algorithm – Failure of nodes



## Massive failures have little impact

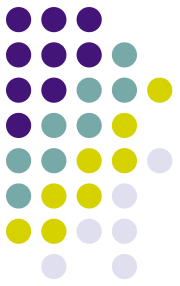




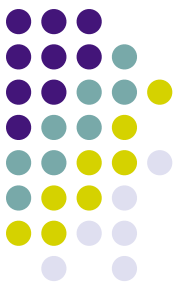
# Overview

- Introduction
- The Chord Algorithm
  - Construction of the Chord ring
  - Localization of nodes
  - Node joins and stabilization
  - Failure/Departure of nodes
- Applications
- Summary
- Questions

# Applications: Chord-based DNS



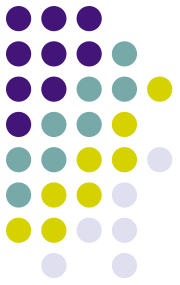
- DNS provides a lookup service  
keys: host names values: IP addresses  
Chord could hash each host name to a key
- Chord-based DNS:
  - no special root servers
  - no manual management of routing information
  - no naming structure
  - can find objects not tied to particular machines



# Summary

- Simple, powerful protocol
- Only operation: map a key to the responsible node
- Each node maintains information about  $O(\log N)$  other nodes
- Lookups via  $O(\log N)$  messages
- Scales well with number of nodes
- Continues to function correctly despite even major changes of the system

# Questions?





**Thanks!**