

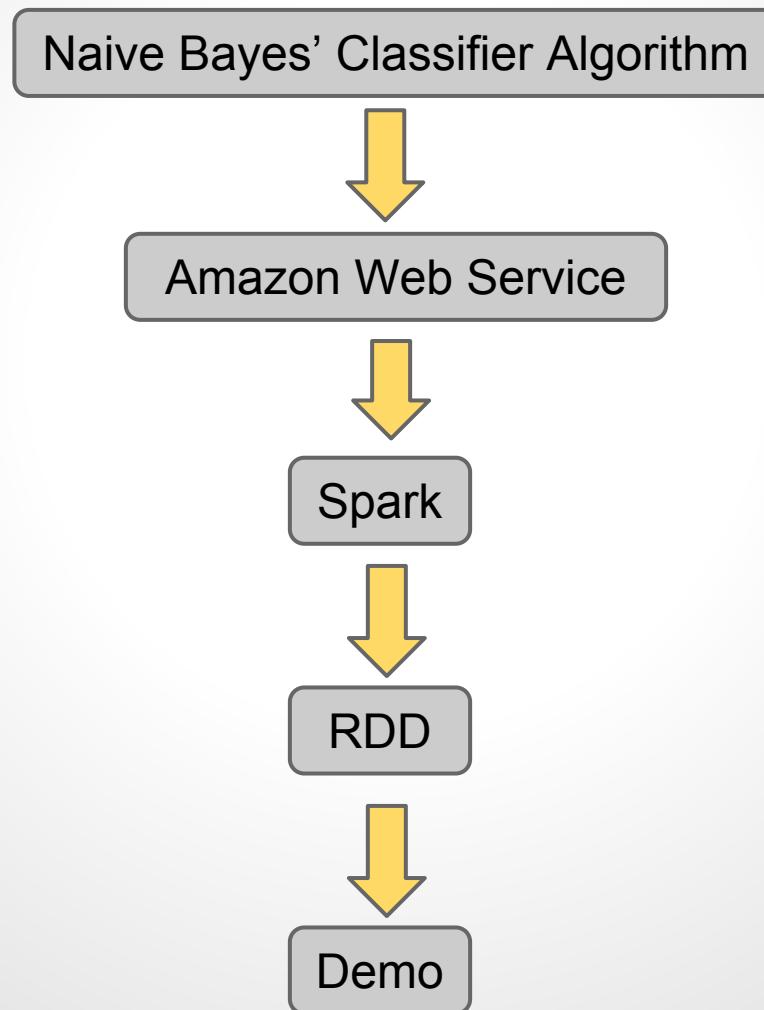
Spark Application on AWS EC2 Cluster

COEN241 - Cloud Computing

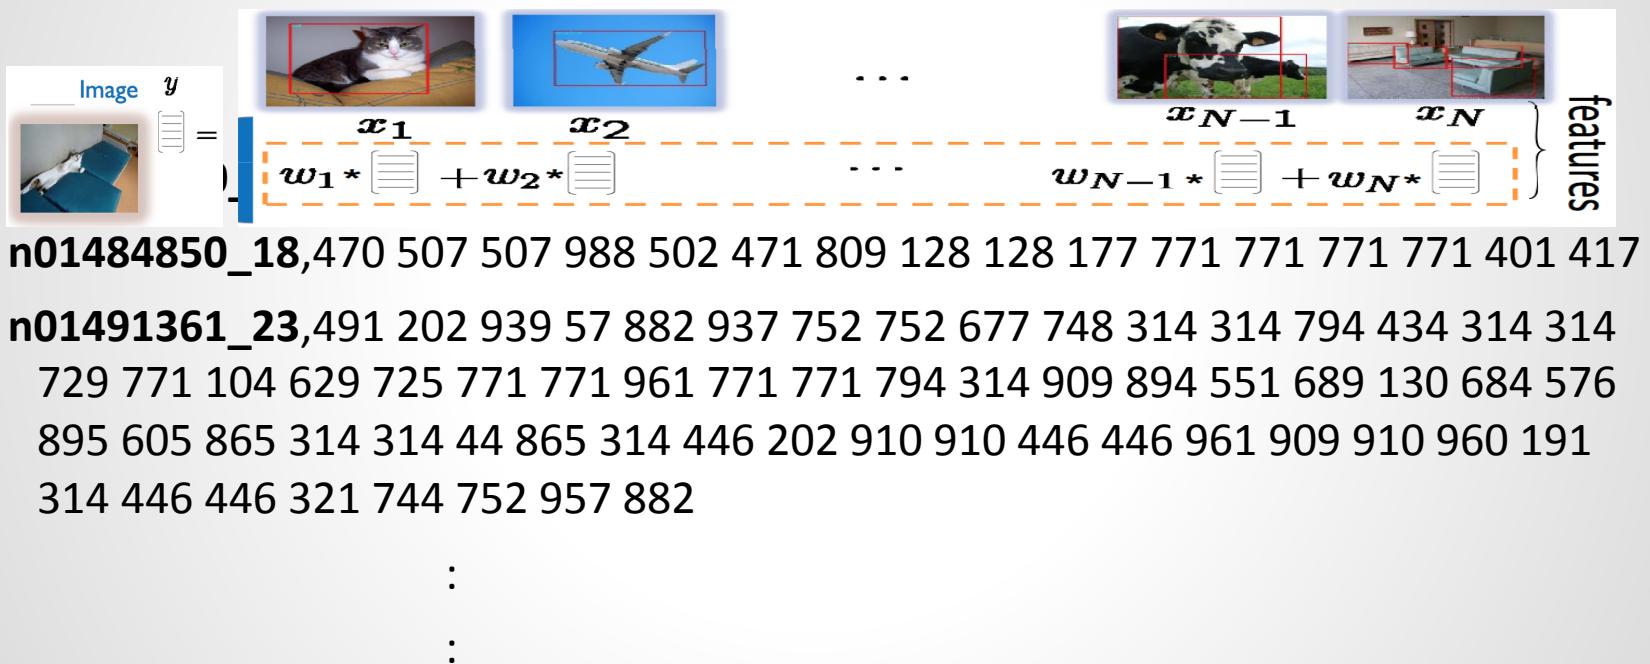
Group 6

Chao-Hsuan Shen, Patrick Loomis, John Geevarghese

Explanation Outline



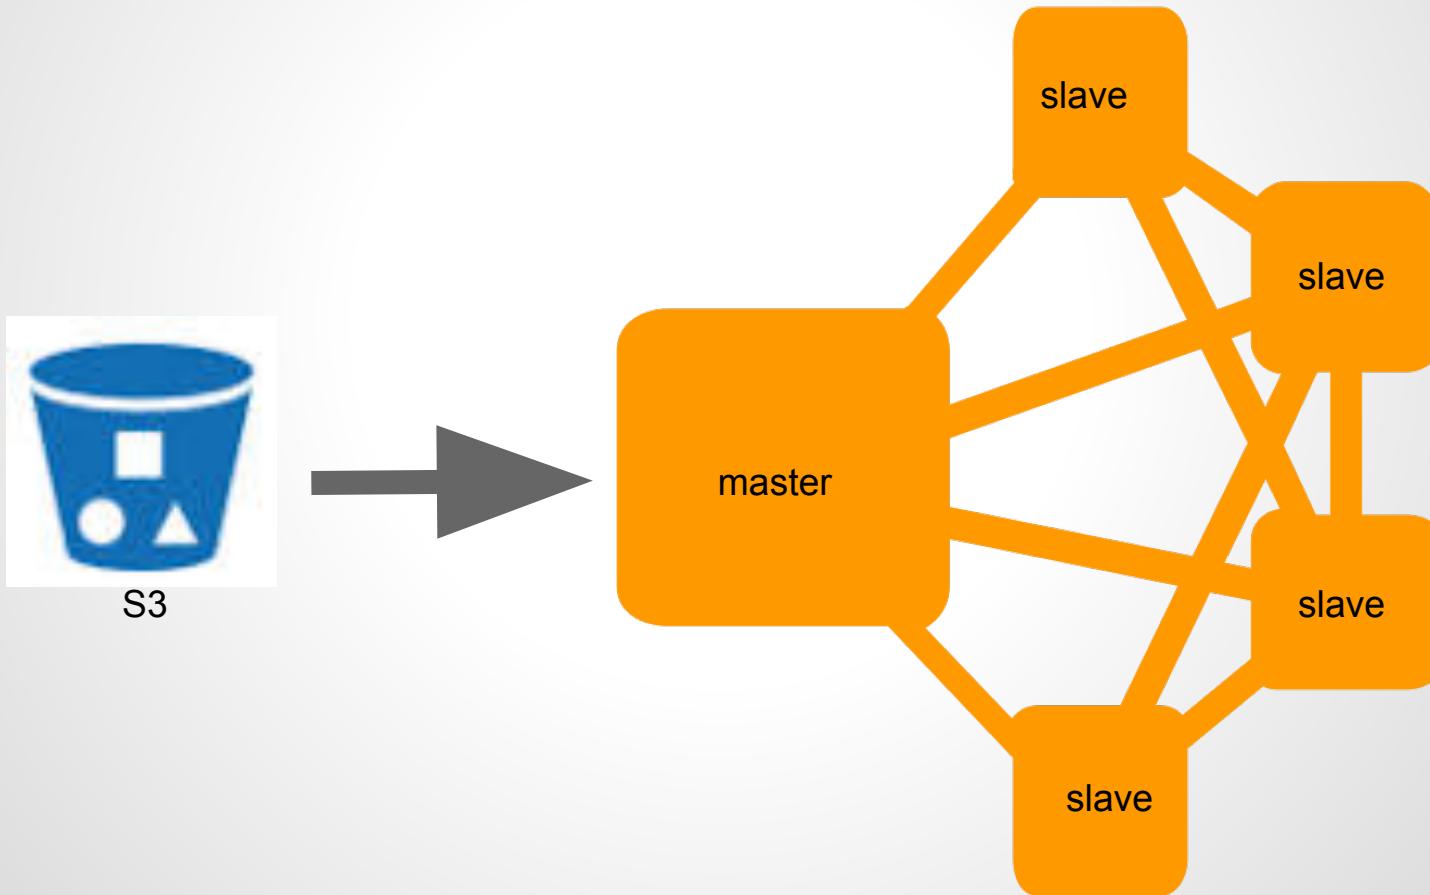
Raw Data of Imagenet



Key Features:

- 1,000 classes/files
- 10 million labeled images, depicting 10,000+ object categories
 - (all done by hand)

AWS - Architecture



AWS - Instance & S3

Instance Type	vCPU	Memory (GiB)	Storage (GB)	Networking Performance	Physical Processor	Clock Speed (GHz)
t2.micro	1	1	EBS Only	Low to Moderate	Intel Xeon family	2.5
t2.small	1	2	EBS Only	Low to Moderate	Intel Xeon family	2.5
t2.medium	2	4	EBS Only	Low to Moderate	Intel Xeon family	2.5
m3.medium	1	3.75	1 x 4 SSD	Moderate	Intel Xeon E5-2670 v2*	2.5
m3.large	2	7.5	1 x 32 SSD	Moderate	Intel Xeon E5-2670 v2*	2.5
m3.xlarge	4	15	2 x 40 SSD	High	Intel Xeon E5-2670 v2*	2.5
m3.2xlarge	8	30	2 x 80 SSD	High	Intel Xeon E5-2670 v2*	2.5

The 'Memory (GiB)' column is highlighted with an orange border.

Services ▾
 EC2
 S3
Edit ▾

Upload
Create Folder
Actions ▾

All Buckets / lucas-coen241

	Name	Storage Class	Size
<input checked="" type="checkbox"/>	1_GB_Train_and_Test_data.tsv	Standard	1 GB
<input type="checkbox"/>	200_MB_Test_dat_tab.tsv	Standard	211.1 MB
<input type="checkbox"/>	550_MB_Train_and_Test_data.tsv	Standard	528.8 MB
<input type="checkbox"/>	FullText.tsv	Standard	3.3 GB

AWS- Cluster

- Stand alone cluster

```
./sbin/start-master.sh
```

```
./bin/spark-class org.apache.spark.deploy.worker.Worker spark://IP:PORT
```

- EC2 cluster

```
./spark-ec2 -k <key pair> -i <key file path> launch <cluster-name>
```

```
./spark-ec2 -k <key pair> -i <key file path> login <cluster-name>
```

```
./spark-ec2 -k <key pair> -i <key file path> start|stop <cluster-name>
```

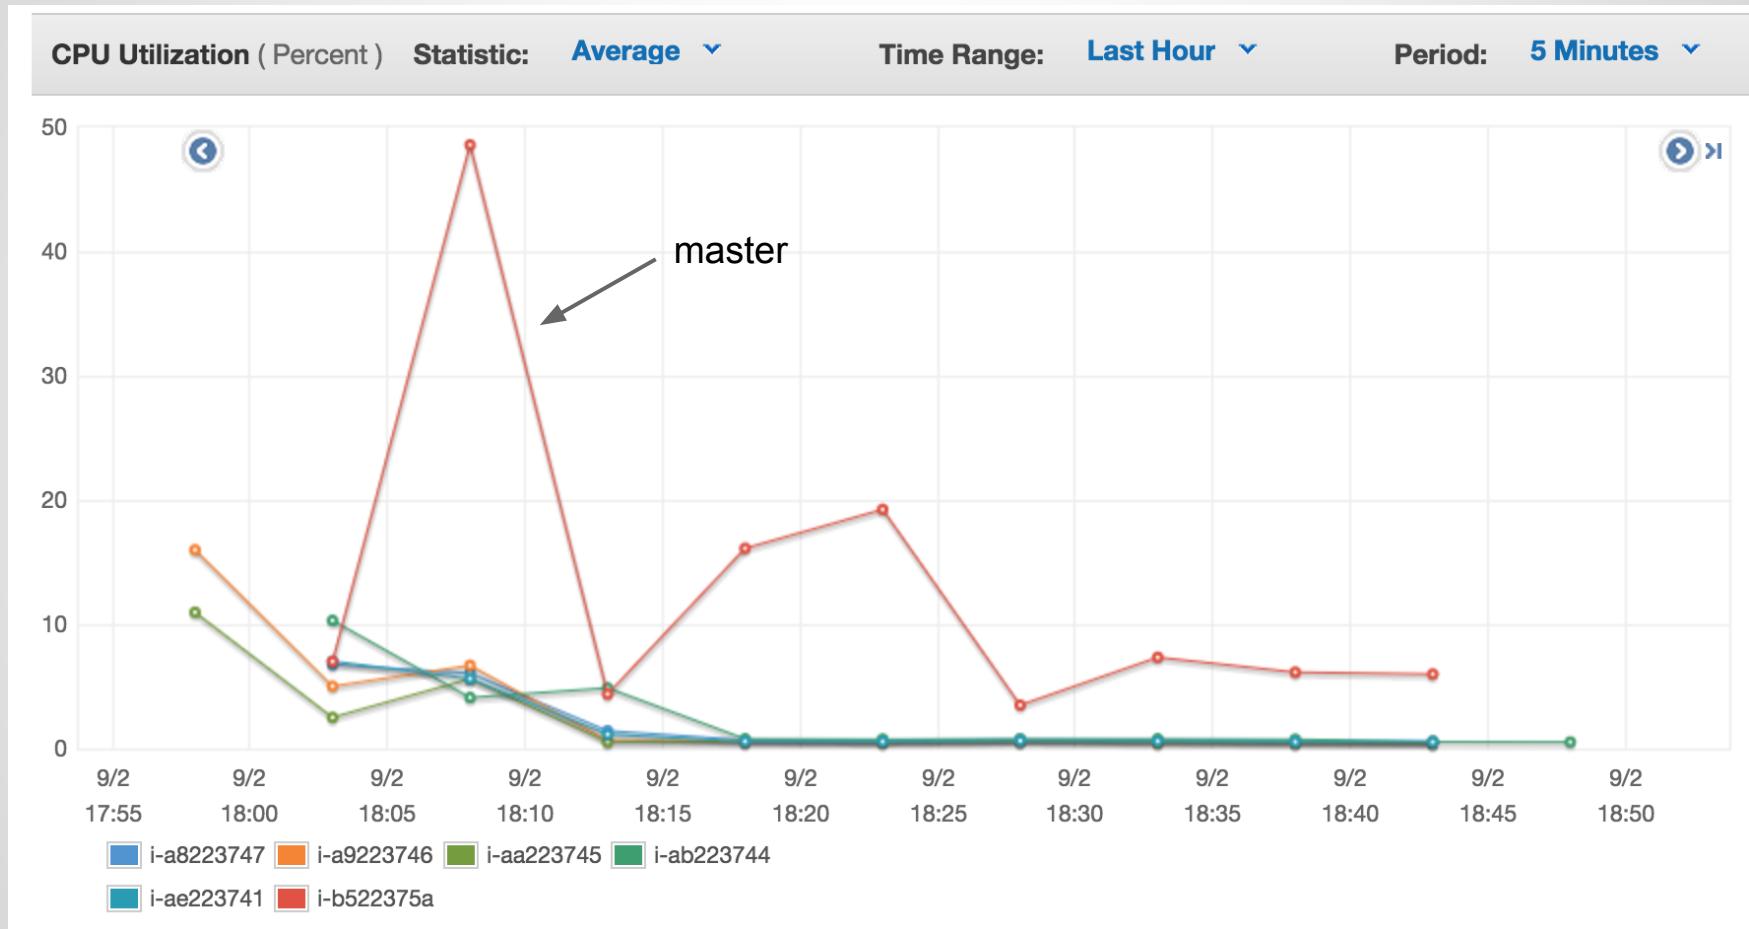
- on YARN && Mesos

- Probably better idea (buggy scripts..)
- for existing distributed data on existing cluster
- Amazon Elastic MapReduce

AWS- Cluster set-up roadmap

1. check identity (ASW key-pair)
2. set-up security group
3. Launch slaves
4. Launch master
5. set-up keyless ssh (rsa)
6. install Apache Mesos (cluster manager)
7. install spark, shark, hdfs, tachyon, ganglia on cluster
8. running setup.sh on master node

AWS- CloudWatch



CPU, Disk Read, Disk Write, Network in, Network out

AWS- Obstacles Faced

- Keyless SSH setup

“Failed to SSH to remote host ec2-54-213-154-133.us-west-2.compute.amazonaws.com.”

When Instance say it is running, it is **not actually fully running...**

- Old version instance vs new version

- m1.large → Good^^
- m3.large → Problems...@.@

- S3 v.s HDFS

- S3 is freaking slow to load big file into cluster. after 40 min loading data from S3... I quit

- Iaas v.s Paas

- update software yourself → older version of AWS CLI..==

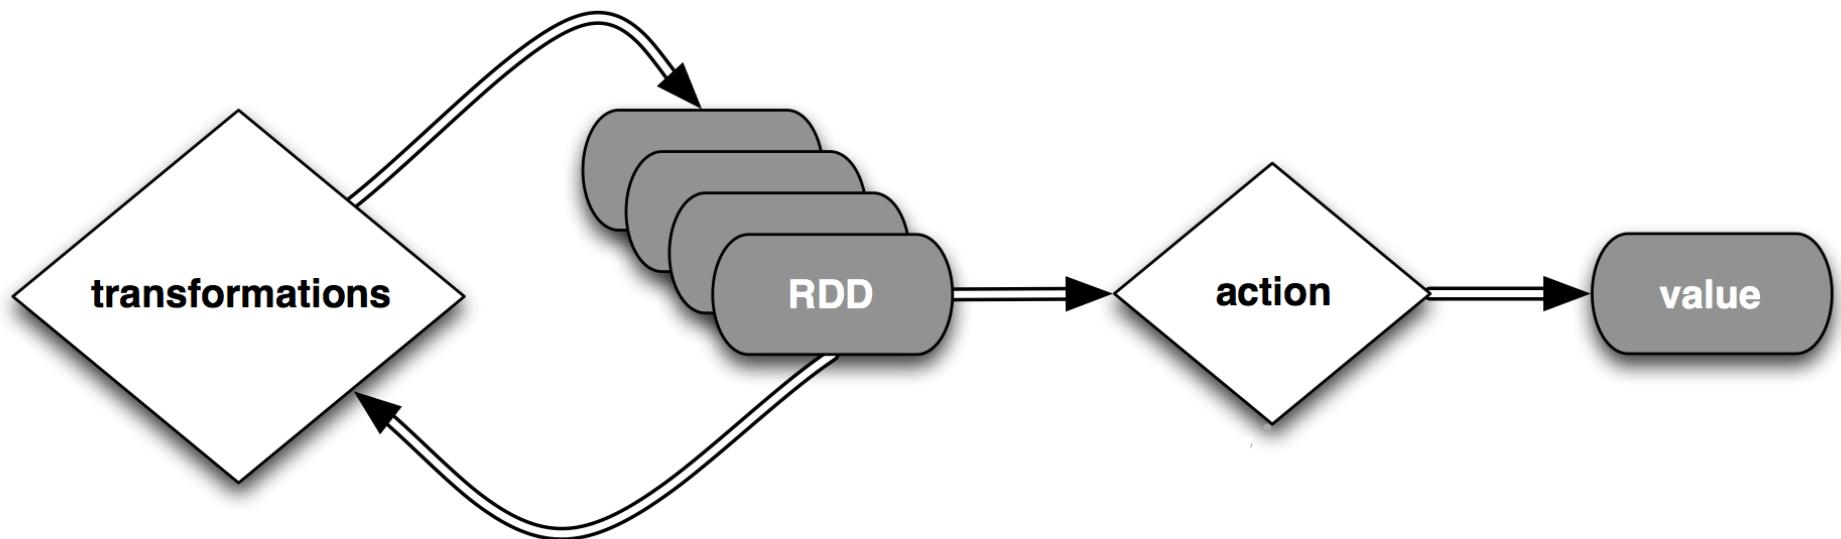
Spark - Scala & Akka

- Functional Programming + Object-Oriented
 - 1. every value is an **object** and every operation is a **method-call**
 - 2. first-class functions
 - 3. a library with **efficient immutable** data structures
- functional programming w/out fully OO

$\text{fn}_2(*\text{args}_2, \text{fn}_1(*\text{args}_1, \text{fn}_0(*\text{args}_0)))$

- **Akka actor** as backbone of Spark

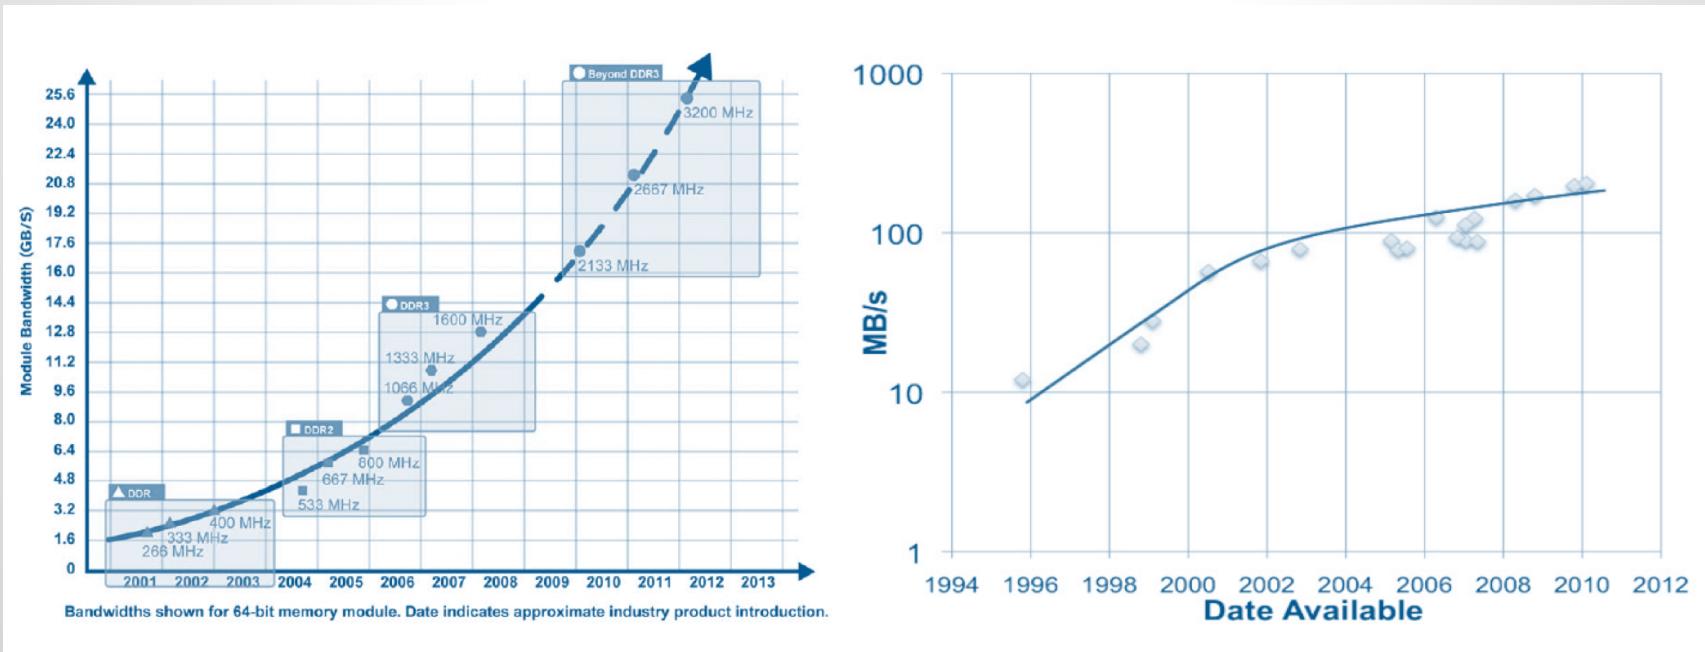
Spark- RDD



Why in-memory processing?

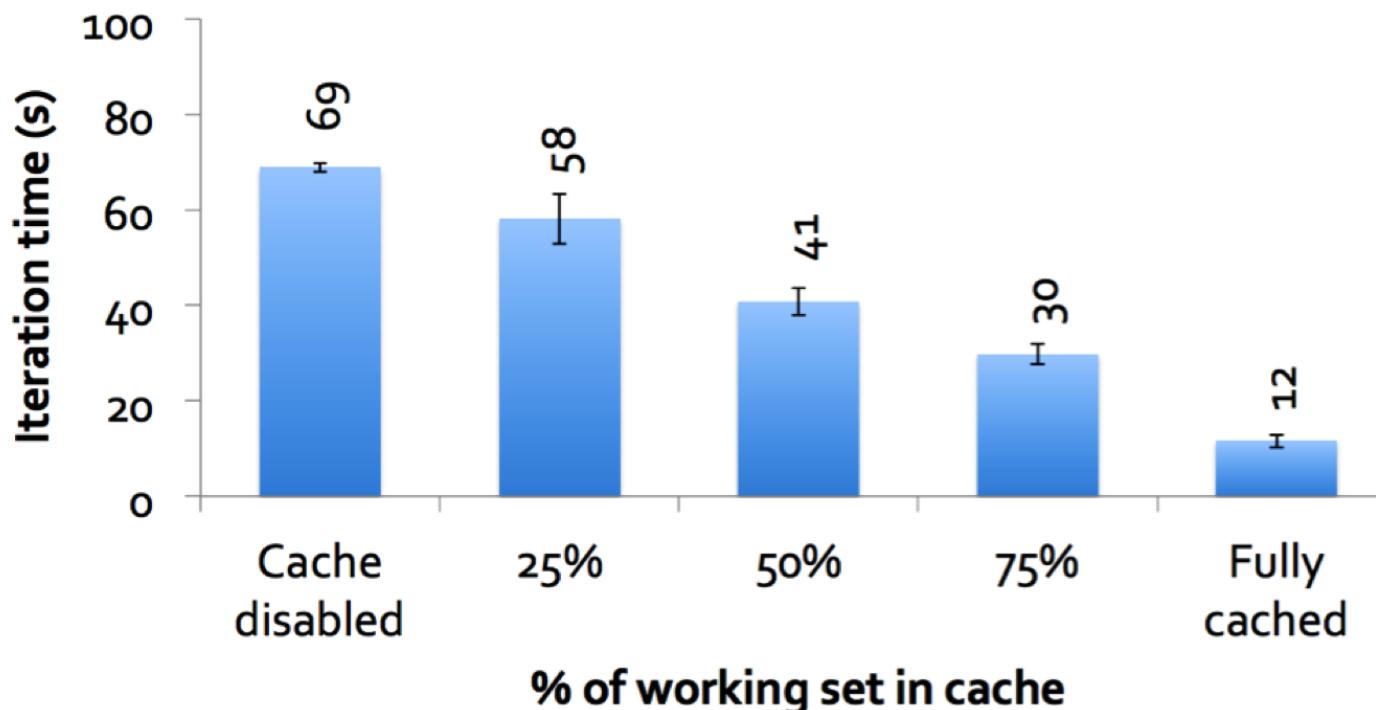
Support **batch**, **streaming**, and **interactive** computations..... in a unified framework.

Memory Throughput matters

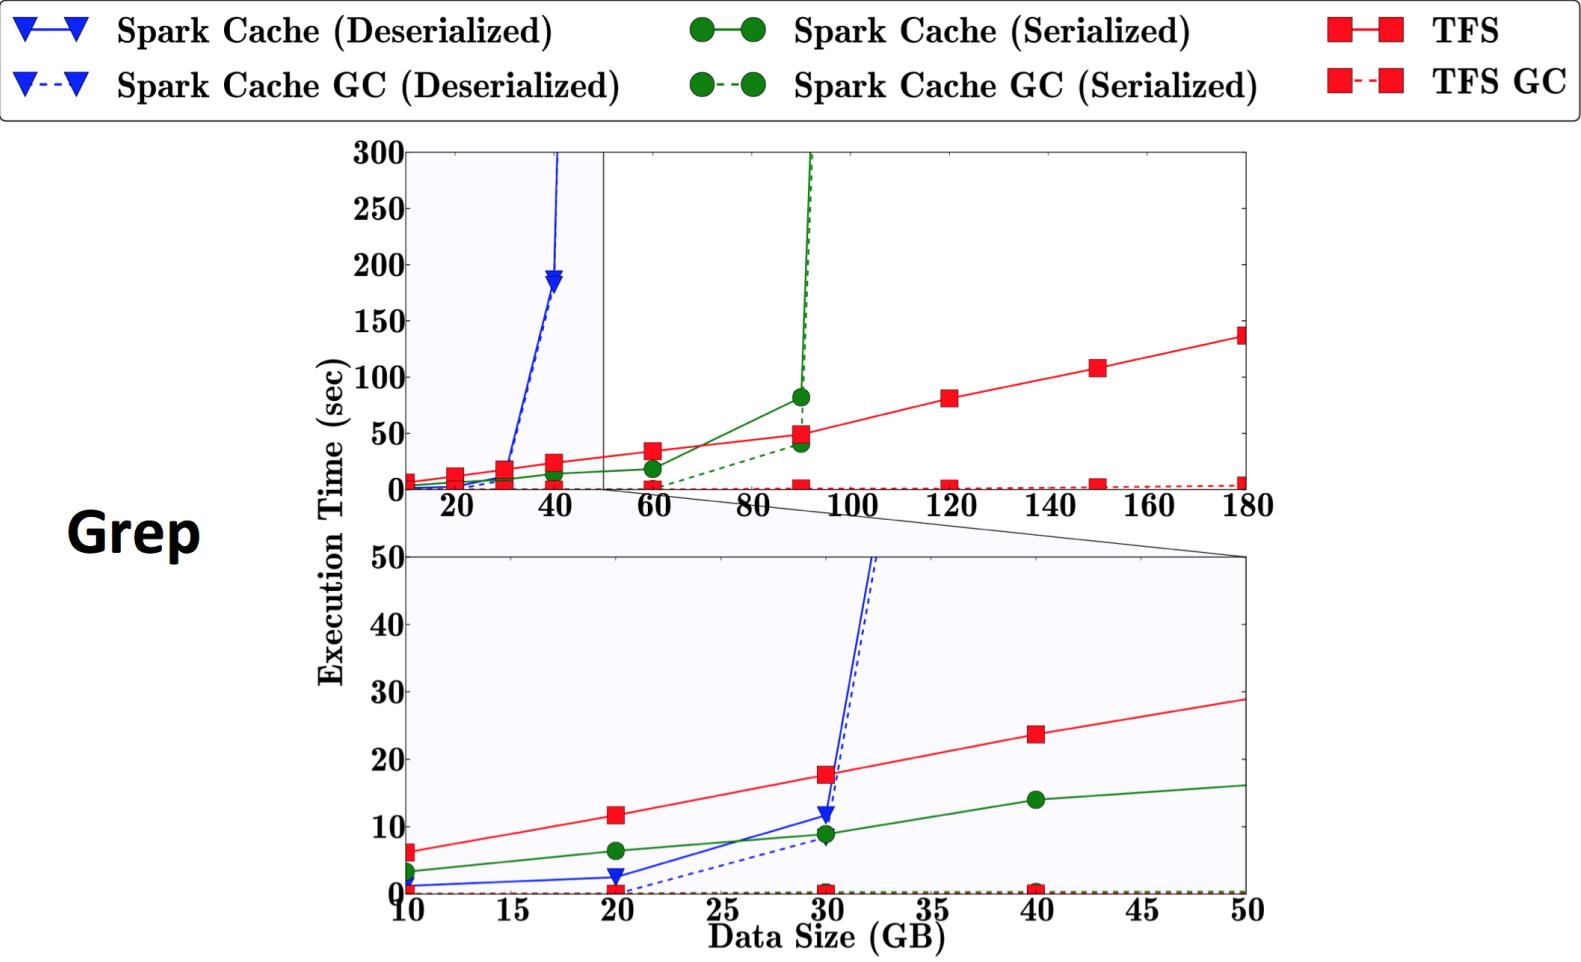


Spark- Cache

Behavior with Less RAM



Spark- JVM vs Tachyon



Spark- Monitoring Cluster



Spark Master at spark://ec2-54-85-149-164.compute-1.amazonaws.com:7077

URL: spark://ec2-54-85-149-164.compute-1.amazonaws.com:7077

Workers: 5

Cores: 10 Total, 10 Used

Memory: 31.4 GB Total, 30.0 GB Used

Applications: 1 Running, 9 Completed

Drivers: 0 Running, 0 Completed

Status: ALIVE

Workers

ID	Address	State	Cores	Memory
worker-20140902200238-ip-172-31-15-122.ec2.internal-52597	ip-172-31-15-122.ec2.internal:52597	ALIVE	2 (2 Used)	6.3 GB (6.0 GB Used)
worker-20140902200238-ip-172-31-15-123.ec2.internal-56830	ip-172-31-15-123.ec2.internal:56830	ALIVE	2 (2 Used)	6.3 GB (6.0 GB Used)
worker-20140902200238-ip-172-31-15-124.ec2.internal-58044	ip-172-31-15-124.ec2.internal:58044	ALIVE	2 (2 Used)	6.3 GB (6.0 GB Used)
worker-20140902200238-ip-172-31-15-125.ec2.internal-50410	ip-172-31-15-125.ec2.internal:50410	ALIVE	2 (2 Used)	6.3 GB (6.0 GB Used)
worker-20140902200238-ip-172-31-15-126.ec2.internal-54495	ip-172-31-15-126.ec2.internal:54495	ALIVE	2 (2 Used)	6.3 GB (6.0 GB Used)

Running Applications

ID	Name	Cores	Memory per Node	Submitted Time	User	State	Duration
app-20140902212320-0009	NaiveBayes Application	10	6.0 GB	2014/09/02 21:23:20	root	RUNNING	3.3 min

Completed Applications

ID	Name	Cores	Memory per Node	Submitted Time	User	State	Duration
app-20140902211703-0008	NaiveBayes Application	10	6.0 GB	2014/09/02 21:17:03	root	FINISHED	3.4 min
app-20140902211008-0007	NaiveBayes Application	10	6.0 GB	2014/09/02 21:10:08	root	FINISHED	3.5 min
app-20140902210625-0006	NaiveBayes Application	10	6.0 GB	2014/09/02 21:06:25	root	FINISHED	3.4 min
app-20140902210220-0005	NaiveBayes Application	10	6.0 GB	2014/09/02 21:02:20	root	FINISHED	3.5 min
app-20140902205423-0004	NaiveBayes Application	10	6.0 GB	2014/09/02 20:54:23	root	FINISHED	3.5 min

Spark- Monitoring running App

Spark Stages

Total Duration: 6.3 min

Scheduling Mode: FIFO

Active Stages: 1

Completed Stages: 6

Failed Stages: 0

Active Stages (1)

Stage Id	Description	Submitted	Duration	Tasks: Succeeded/Total	Shuffle Read	Shuffle Write
7	(kill) count at naive.scala:77	2014/09/02 21:24:27	5.1 min	0/2		

Completed Stages (6)

Stage Id	Description	Submitted	Duration	Tasks: Succeeded/Total	Shuffle Read	Shuffle Write
5	collect at NaiveBayes.scala:96	2014/09/02 21:24:22	1 s	2/2	1932.6 KB	
6	combineByKey at NaiveBayes.scala:91	2014/09/02 21:23:49	33 s	2/2		3.8 MB
4	collect at naive.scala:28	2014/09/02 21:23:33	16 s	2/2		
2	collect at naive.scala:21	2014/09/02 21:23:32	1 s	2/2	8.3 KB	
0	apply at Option.scala:120	2014/09/02 21:23:31	1 s	1/1	4.3 KB	
1	distinct at naive.scala:19	2014/09/02 21:23:20	10 s	2/2		8.3 KB

Failed Stages (0)

Stage Id	Description	Submitted	Duration	Tasks: Succeeded/Total	Shuffle Read	Shuffle Write	Failure Reason

Mahout vs MLlib

Future

Mahout had fairly straight forward operations

- Sequence file → Sparse Matrix → TF-IDF → train model

Spark MLlib did not provide all of the functions needed to classify

- Text classification was initially implemented in MLlib, but not the stable release of MLlib did not include any of the functions needed
- TFHashing, Word2Vec, and a few other functions are being implemented in an upcoming stable release - *Spark 1.1.0*

If we had more time we could contribute our own implementation of text classification using multinomial naive bayes':



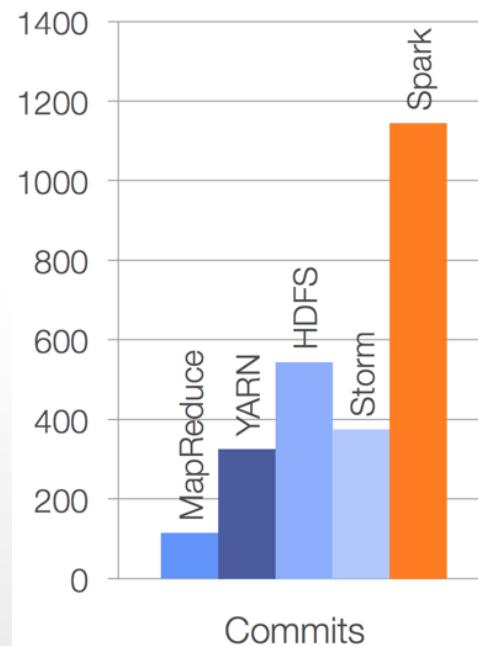
Mahout vs MLlib - Ecosystem

Mahout is moving away from its implementation of MapReduce, and moving towards DSL for linear algebraic operations

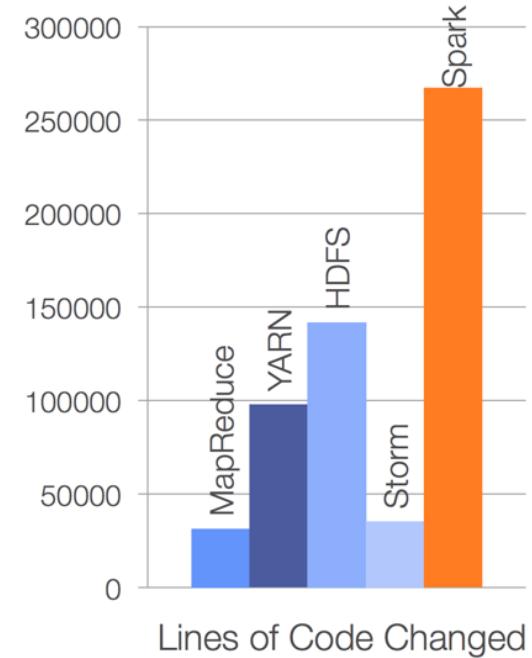
- which are automatically optimized and executed in parallel on Spark
- Unfortunately, NB is still in development for Spark

Contributions to Spark are rising dramatically:

	June 2013	July 2014
Total Contributions	68	255
Companies Contributing	17	50
Total lines of code	63,000	175,000

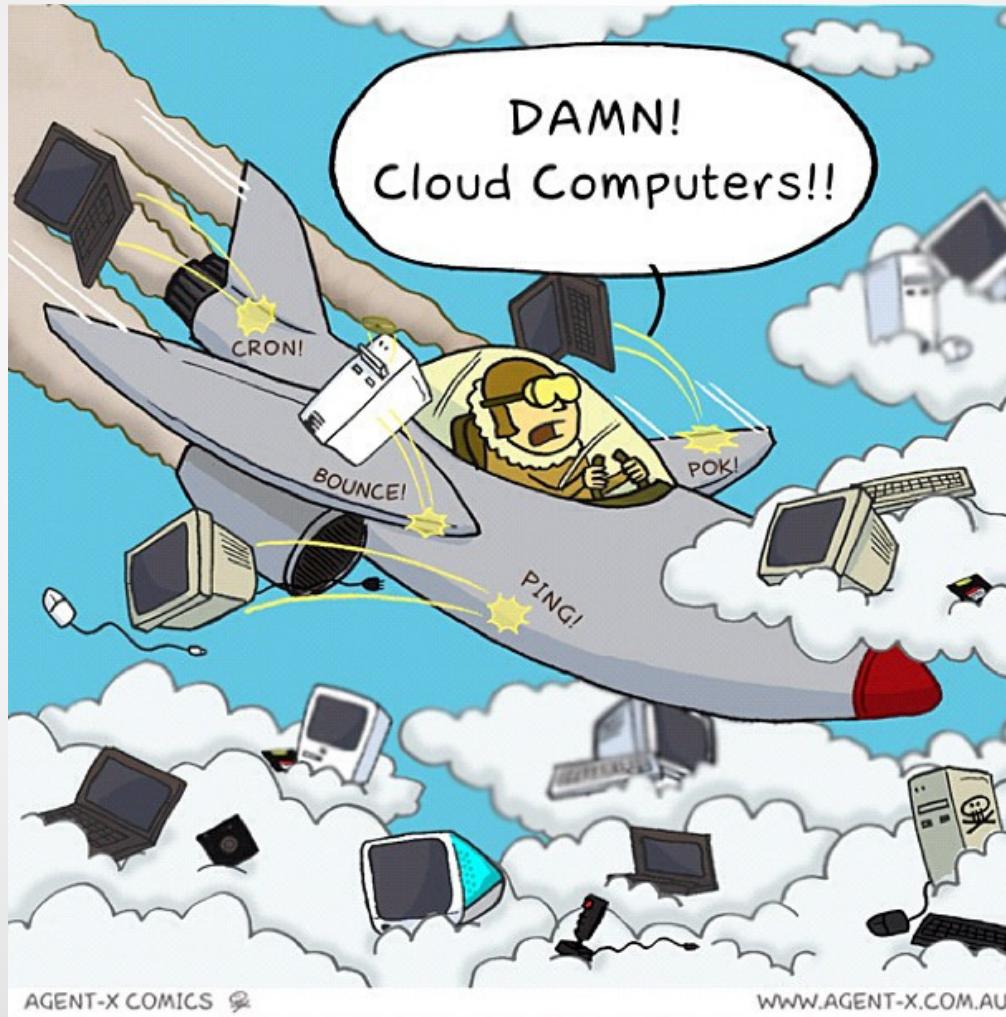


Activity in past 6 months



DEMO

Fin



Naive Bayes' Classifier (1)

Two things of note:

1. Conditional Probability

- i. What is the probability that something will happen, given that something else has already happened?
- ii. Example:

2. Bayes' Rule

- i. Predict an outcome given *multiple evidence*
- ii. ‘uncouple’ multiple pieces of evidence and treat each piece of evidence as independent = *naive Bayes*
- iii. Multiple *prior* so that we give high probability to more common outcomes, and low probabilities to unlikely outcomes. These are also called *base rates* and they are a way to scale our predicted probabilities

Naive Bayes' Classifier (2) - TF-IDF

Convert word counts of a document using the TF-IDF transformation before applying the learning algorithm. The resulting formula gives a weighted importance of word

- f = word frequency
- D = number of documents
- df = number of documents containing word under consideration

$$TFIDF(word) = \log(f + 1) \times \log\left(\frac{D}{df}\right).$$

Naive Bayes' Classifier (3)

Naive Bayes Multinomial

- Class label denoted by C; n is size of vocabulary
- Class Prior, $\Pr(c) = \# \text{ of documents belonging to class } c \text{ divided by the total } \# \text{ of documents}$
- Likelihood, $\Pr(x|c) = \text{Probability of obtaining a document like } x \text{ in class } c$
- Predictor Prior (Evidence), $\Pr(x) = \text{count of a feature over the total count of features}$
- MNB assigns a test document x to the class that has the highest probability $\Pr(c | x)$

$$P(c | x) = \frac{P(x | c)P(c)}{P(x)}$$

Likelihood Class Prior Probability
↓ ↑
Posterior Probability Predictor Prior Probability

$$P(c | X) = P(x_1 | c) \times P(x_2 | c) \times \dots \times P(x_n | c) \times P(c)$$